

Google BigQuery Lab

1	BIGQUERY BACKGROUND.....	1
2	LAB SETUP	1
3	USING THE COUNT FUNCTION	4
4	USING COUNT WITH NULL RECORDS	6
5	RETRIEVING A DISTINCT SET OF VALUES	7
6	BASIC USE OF WHERE	8
7	USING WHERE TO FILTER ON NUMERICAL COLUMNS WITH OPERATORS	8
8	USING WHERE TO FILTER ON STRING COLUMNS WITH OPERATORS.....	11
9	USING WHERE ON TIMESTAMP COLUMNS	12
10	USING WHERE TO FILTER ON NULL / NOT NULL.....	14
11	INTRODUCING GROUP BY.....	15
12	USING GROUP BY WITH AGGREGATE FUNCTIONS	17
13	PERFORMING JOINS	18
14	LOOKER STUDIO BACKGROUND	18

1 BigQuery background

General overview of BigQuery

<https://cloud.google.com/bigquery/docs/introduction>

Overview of BigQuery Data Storage

https://cloud.google.com/bigquery/docs/storage_overview

Overview of BigQuery analytics

<https://cloud.google.com/bigquery/docs/query-overview>


2 Lab setup

You can use the BigQuery sandbox to explore limited BigQuery capabilities without providing a credit card or creating a billing account for your project. If you already created a billing account, you can still use BigQuery at no cost in the free usage tier.

<https://console.cloud.google.com/bigquery>

Once you are at the main Big Query page, you can create a new project with the name: My First BigQuery Project


New Project

 You have 12 projects remaining in your quota. Request an increase or delete projects. [Learn more](#)

[MANAGE QUOTAS](#)

Project name *
My First BigQuery Project

Project ID: upbeat-glow-393712. It cannot be changed later. [EDIT](#)

Location *
 No organization [BROWSE](#)

Parent organization or folder

[CREATE](#) [CANCEL](#)


https://cloud.google.com/bigquery/docs/quickstarts/query-public-dataset-console#open_a_public_dataset

Once done, click Add in Explorer, and select Public DataSets.


SANDBOX Set up billing to upgrade to the

Explorer [+ ADD](#) [I<](#)


Add

 Google Cloud


Google storage service

 Salesforce Data Cloud


Data published from Salesforce platform

 Amazon S3 - Data Transfer

Amazon object storage service, via the Data Transfer Service

 Azure Blob Storage (and Azure Data Lake Storage Gen2) - Data Transfer

Microsoft object storage service (and data lake storage service), via the Data Transfer Service

 Public Datasets

BigQuery public datasets from the Google Cloud Public Dataset Program

Select the first dataset in this list and click View Data Set.

Marketplace

Marketplace > Data

Filter Type to filter

235 results

Category


Analytics (43)

Big data (32)

Databases (6)


Machine learning (6)

Generative AI (2)

 Cloud-to-Ground Lightning Strikes

NOAA

Aggregated lightning strike data from 1987 to 2018

 **Cloud-to-Ground Lightning Strikes**

[NOAA](#)

Aggregated lightning strike data from 1987 to 2018

[VIEW DATASET](#)

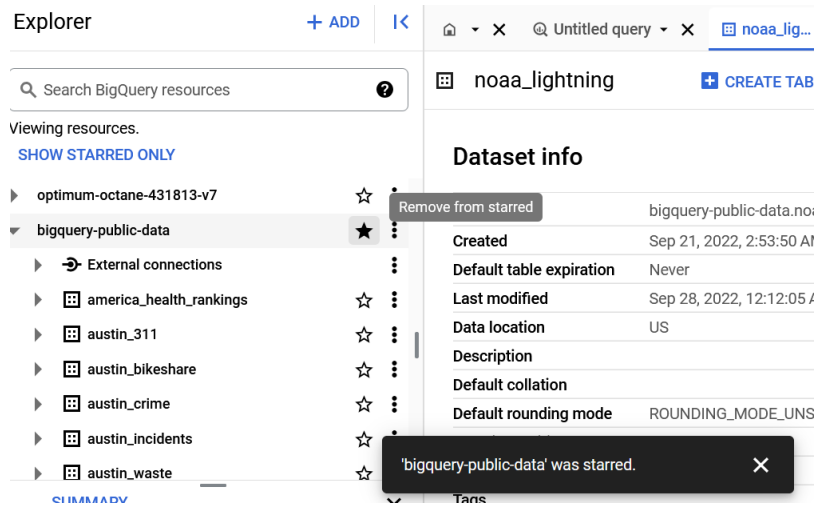
[Click to view dataset](#)

[OVERVIEW](#)

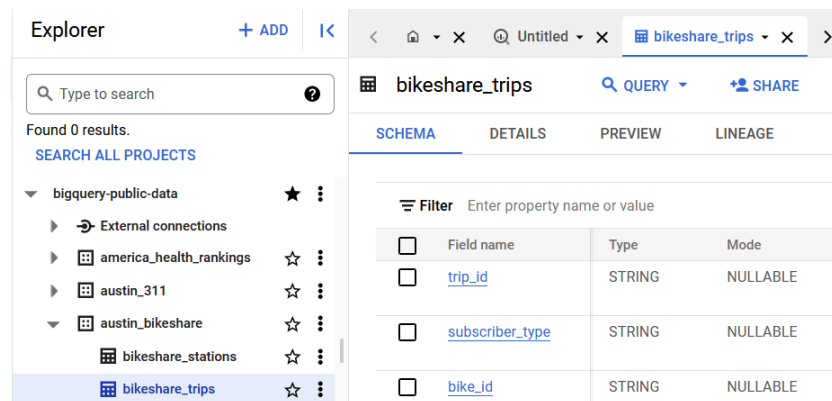
[SAMPLES](#)

[RELATED PRODUCTS](#)

This will add the `bigquery-public-data` option in the left pane, which you can proceed to star.



Opening up any one of the datasets allows us to select a table which opens up a tab in the details pane, which we can then explore in more detail by selecting the appropriate tabs (for e.g. Schema, Details, Preview and Lineage for a normal table).



Notice that for some of these tables, the number of rows are in the millions and tens of millions. This exceeds the capacity of small scale data analytics tools such as Excel:

<https://support.microsoft.com/en-gb/office/excel-specifications-and-limits-1672b34d-7043-467e-8e27-269d656771c3>

Even popular relational databases such as MySQL have hard limits on their table column count and row size

<https://dev.mysql.com/doc/refman/8.0/en/column-count-limit.html>

The number of databases / tables are not limited by the system, but by the underlying file storage system and there are limits with modern SAN storage systems:

<https://dev.mysql.com/doc/refman/8.0/en/database-count-limit.html>

We will demonstrate some basic analytics that we can perform on a table using query statements (or Data Query Language (DQL) statements) written in GoogleSQL. The statements operate by scanning one or more specified tables and returns the computed result rows.

<https://cloud.google.com/bigquery/docs/introduction-sql>

3 Using the Count function

Search for this table in the `bigquery-public-data`, select and star it and view it in the details pane and explore it in more detail

`usa_names.usa_1910_2013`



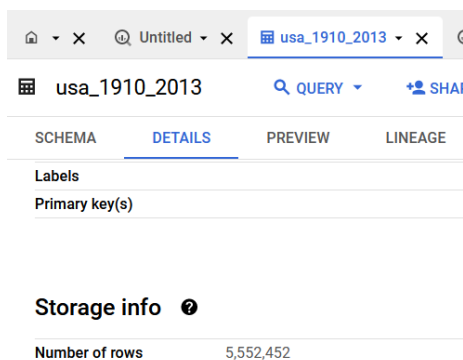
The most basic query (which you can type in a new query tab in the details pane) is to see the first 1000 (or whatever number you desire) number of rows / records in the table.

```
SELECT * FROM `bigquery-public-data.usa_names.usa_1910_2013` LIMIT 1000
```

You can count the total number of row / records with:

```
SELECT
  COUNT(*) AS num_records
FROM
  `bigquery-public-data.usa_names.usa_1910_2013`
```

Notice that this count tallies with the info about the table:



To get the count of the number of rows with the column name (which will be the same as the total number of rows in the table, since all rows have this column):

```
SELECT
  COUNT(name) AS cnt
FROM
  `bigquery-public-data.usa_names.usa_1910_2013`
```

To see how many distinct (unique) years there are:

```
SELECT
  COUNT(DISTINCT year) AS distinct_year_count
FROM
  `bigquery-public-data.usa_names.usa_1910_2013`
```

The result reflects the 104 distinct and different years between 1910 and 2013

To see how many distinct (unique) states there are:

```
SELECT
  COUNT(DISTINCT state) AS distinct_state_count
FROM
  `bigquery-public-data.usa_names.usa_1910_2013`
```

There are 51 states (50 states + Washington DC)

To see how many distinct (unique) genders there are:

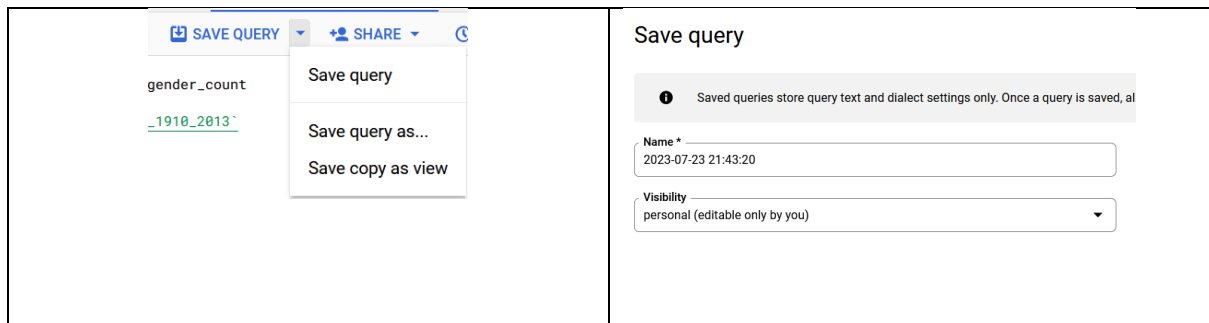
```
SELECT
  COUNT(DISTINCT gender) AS distinct_gender_count
FROM
  `bigquery-public-data.usa_names.usa_1910_2013`
```

We can combine all these separate queries into a single larger query if we wish:

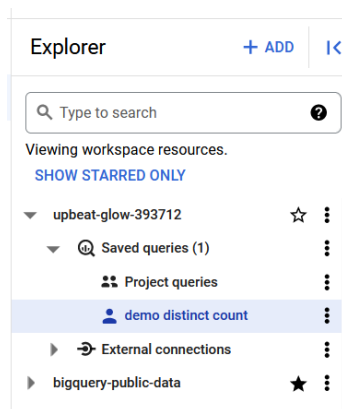
```
SELECT
  COUNT(DISTINCT gender) AS distinct_gender_count,
  COUNT(DISTINCT year) AS distinct_year_count,
  COUNT(DISTINCT state) AS distinct_state_count,
  COUNT(DISTINCT name) AS distinct_name_count,
  COUNT(*) AS num_records,
  COUNT(name) AS cnt
FROM
  `bigquery-public-data.usa_names.usa_1910_2013`
```

You can save any of these queries if you wish for future reference / reuse in the same project or to share with others:

--	--



The saved query should be visible in the Explorer Pane:



4 Using Count with Null records

Often, large datasets with have null values in some of their columns due to the various issues encountered when gathering the data.

Look for `new_york_mv_collisions.nypd_mv_collisions` in `bigquery-public-data` and star it in the Explorer Pane.

You can verify the number of rows with:

```
SELECT
  COUNT(*) AS num_rows
FROM
  `bigquery-public-data.new_york_mv_collisions.nypd_mv_collisions`
```

However, if you check again with:

```
SELECT
  COUNT(contributing_factor_vehicle_1) AS
  contributing_factor_vehicle_1_count
FROM
  `bigquery-public-data.new_york_mv_collisions.nypd_mv_collisions`
```

You will notice the count is slightly less than the total number of rows, indicating the presence of NULL in some of the rows of this column.

This is even more evident in the next query:

```
SELECT
  COUNT(contributing_factor_vehicle_2) AS
contributing_factor_vehicle_2_count
FROM `bigquery-public-
data.new_york_mv_collisions.nypd_mv_collisions`
```

This can also be seen in the preview tab of the details pane for this table.

We can also check the number of null values in the other columns with:

```
SELECT
  COUNT(borough) AS borough_count,
  COUNT(contributing_factor_vehicle_3) AS
contributing_factor_vehicle_3_count,
FROM `bigquery-public-data.new_york_mv_collisions.nypd_mv_collisions`
```

5 Retrieving a distinct set of values

You can retrieve a distinct set of values from a column and order the results as well using the ORDER BY clause

For example, if we wanted to return the distinct set of years, starting with the most recent, from the USA names data set we would run this query:

```
SELECT
  DISTINCT year
FROM `bigquery-public-data.usa_names.usa_1910_current`
ORDER BY
  year DESC
```

Or if we wanted to return a list of the first 50 names sorted in alphabetical order we could execute this query:

```
SELECT
  DISTINCT name
FROM `bigquery-public-data.usa_names.usa_1910_current`
ORDER BY
  name ASC
LIMIT
  50
```

If we wanted to see the last 100 names we could execute:

```
SELECT
```

```
DISTINCT name
FROM
  `bigquery-public-data.usa_names.usa_1910_current`
ORDER BY
  name DESC
LIMIT
  100
```

You can select from the distinct values from a different number of columns at the same time and order the results based on one or more of these columns simultaneously:

```
SELECT
  DISTINCT
    borough,
    major_category,
    minor_category
FROM
  `bigquery-public-data.london_crime.crime_by_lsoa`
ORDER BY
  borough,
  major_category,
  minor_category
```

6 Basic use of WHERE

The WHERE clause allows you to specify a search condition for the rows returned by a query. The `search_condition` is a combination of one or more expressions using the logical operator AND, OR and NOT.

We can select conditions involving single columns:

```
SELECT * FROM `bigquery-public-data.usa_names.usa_1910_current`
WHERE state = 'FL'
```

```
SELECT * FROM `bigquery-public-data.usa_names.usa_1910_current`
WHERE gender = 'M'
```

Or we can have conditions involving the combination of multiple columns

```
SELECT * FROM `bigquery-public-data.usa_names.usa_1910_current`
WHERE state = 'FL' AND gender = 'M' AND year = 2000 ORDER BY
number DESC LIMIT 100
```

7 Using WHERE to filter on numerical columns with operators

We work with the table: `new_york_citibike.citibike_trips`

We can also filter on a numerical column using various combination of operators:

```
SELECT
  count(*) as num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration = 432
```

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration != 432
```

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration < 300
```

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration <= 300
```

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration > 24*60*60
```

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration >= 24*60*60
```

We can also combine multiple conditions using AND or OR operators.

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration >= 5*60*60 AND tripduration <= 9*60*60
```

Notice the use of parenthesis around the two statements separated by OR. You need these brackets when chaining together AND and OR conditions within the same clause.

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  (tripduration >= 1*60*60 AND tripduration <= 3*60*60)
  OR
  (tripduration >= 5*60*60 AND tripduration <= 8*60*60)
```

The IN statement can be used to provide a list within a condition. In this example we are counting the records where the tripduration is in the list (60,120).

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration IN (60,120)
```

The above query is the same as saying the tripduration is 60 or the trip duration is 120

Another similar example:

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration = 60 or tripduration = 120 or tripduration = 180 or
  tripduration = 240
```

can be replaced with a query using an IN statement that is easier to write/read.

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration IN (60,120,180,240)
```

You can also perform a negation of the above query, for e.g. count the number of records where the trip duration was not one of (60,120,180,240).

```
SELECT
  COUNT(*) AS num_bike_rides
FROM
  `bigquery-public-data.new_york_citibike.citibike_trips`
WHERE
  tripduration NOT IN (60,120,180,240)
```

8 Using WHERE to filter on string columns with operators

We will work with the table: `london_crime.crime_by_lsoa`

We can also use the WHERE clause on the contents of text (string) columns for example:

```
SELECT
  COUNT(*) AS num_crimes
FROM
  `bigquery-public-data.london_crime.crime_by_lsoa`
WHERE
  minor_category != "Harassment"
```

```
SELECT
  COUNT(*) AS num_crimes
FROM
  `bigquery-public-data.london_crime.crime_by_lsoa`
WHERE
  minor_category in ("Harassment", "Assault with Injury")
```

```
SELECT
  COUNT(*) AS num_crimes
FROM
  `bigquery-public-data.london_crime.crime_by_lsoa`
WHERE
  minor_category = 'Harassment' or minor_category = 'Assault with Injury'
```

A common technique is to search for a specific pattern within a string column. - You can look for a pattern anywhere in the string by using like '%pattern%' - You can look for a pattern at the end of the string by using like 'pattern' - You can look for a pattern at the start of the string by using like 'pattern%' - The pattern given is any pattern of characters and it's case sensitive.

```
SELECT
```

```
    distinct minor_category
FROM
  `bigquery-public-data.london_crime.crime_by_lsoa`
WHERE
  minor_category like 'Drug%'
```

```
SELECT
  distinct minor_category
FROM
  `bigquery-public-data.london_crime.crime_by_lsoa`
WHERE
  minor_category like '%Drugs'
```

```
SELECT
  distinct minor_category
FROM
  `bigquery-public-data.london_crime.crime_by_lsoa`
WHERE
  minor_category like '%eh%'
```

Often when looking for a pattern, we lower case the column in which we are looking. This way we can just use lower case in our like statement. Google SQL provides many more of these functions that we can support complex query functionality or simplify existing ones.

```
SELECT
  distinct minor_category
FROM
  `bigquery-public-data.london_crime.crime_by_lsoa`
WHERE
  lower(minor_category) like '%motor%'
```

9 Using WHERE on TIMESTAMP columns

We will work with the table: `austin_bikeshare.bikeshare_trips`

A `TIMESTAMP` column will typically have the year, month, day, hour, minute and second. For example 2014-10-26 15:12:00 UTC. A `DATE` column will just have the year, month and day. For example, 2014-10-26. You can change a `TIMESTAMP` into a `DATE` by casting it i.e. `cast(TIMESTAMP AS DATE)`. This would convert 2014-10-26 15:12:00 UTC into 2014-10-26.

```
SELECT
  start_time as start_time_timestamp
FROM
  `bigquery-public-data.austin_bikeshare.bikeshare_trips`
LIMIT
  100
```

```
SELECT
  cast(start_time as date) as start_time_date
FROM
  `bigquery-public-data.austin_bikeshare.bikeshare_trips`
LIMIT
  100
```

```
SELECT
  cast(start_time as date) as start_time_date,
  extract(hour from start_time) as start_time_hour,
  extract(minute from start_time) as start_time_minute
FROM `bigquery-public-data.austin_bikeshare.bikeshare_trips`
LIMIT
  100
```

```
SELECT
  cast(start_time as date) as start_time_date,
  extract(day from start_time) as start_time_day,
  extract(year from start_time) as start_time_year,
  extract(month from start_time) as start_time_month,
  extract(week from start_time) as start_time_week
FROM
  `bigquery-public-data.austin_bikeshare.bikeshare_trips`
LIMIT
  100
```

You can also filter for records after a given date

```
SELECT
  cast(start_time as date) as start_time_date,
  extract(day from start_time) as start_time_day,
  extract(year from start_time) as start_time_year,
  extract(month from start_time) as start_time_month,
  extract(week from start_time) as start_time_week
FROM
  `bigquery-public-data.austin_bikeshare.bikeshare_trips`
WHERE start_time > '2018-10-01'
LIMIT
  100
```

You can obtain records for a specific date

```
SELECT
  cast(start_time as date) as start_time_date,
  extract(day from start_time) as start_time_day,
  extract(year from start_time) as start_time_year,
  extract(month from start_time) as start_time_month,
  extract(week from start_time) as start_time_week
FROM
```

```
`bigquery-public-data.austin_bikeshare.bikeshare_trips`  
WHERE cast(start_time as date) = '2018-10-01'  
LIMIT  
100
```

You can obtain records between two dates:

```
SELECT  
  cast(start_time as date) as start_time_date,  
  extract(day from start_time) as start_time_day,  
  extract(year from start_time) as start_time_year,  
  extract(month from start_time) as start_time_month,  
  extract(week from start_time) as start_time_week  
FROM  
  `bigquery-public-data.austin_bikeshare.bikeshare_trips`  
WHERE start_time >= '2018-09-01' and start_time <= '2018-09-30'  
LIMIT  
100
```

Finally, you can also filter records in a given list of hours:

```
SELECT  
  cast(start_time as date) as start_time_date,  
  extract(hour from start_time) as start_time_hour,  
  extract(minute from start_time) as start_time_minute,  
FROM  
  `bigquery-public-data.austin_bikeshare.bikeshare_trips`  
where extract(hour from start_time) IN (17,18,19,20)  
LIMIT  
100
```

10 Using WHERE To filter on Null / Not Null

We will work with the table: `new_york_mv_collisions.nypd_mv_collisions`

You can check the total number of records in this table from the Details pane, or with:

```
SELECT  
  COUNT(*) AS total_records  
FROM  
  `bigquery-public-data.new_york_mv_collisions.nypd_mv_collisions`
```

To count the total number of records with non-null values in any particular column, we can type:

```
SELECT  
  COUNT(*)
```

```
FROM `bigquery-public-  
data.new_york_mv_collisions.nypd_mv_collisions`  
WHERE  
contributing_factor_vehicle_1 IS NOT NULL
```

Similarly, to count the total number of records with null values in any particular column, we can type:

```
SELECT  
COUNT(*)  
FROM `bigquery-public-  
data.new_york_mv_collisions.nypd_mv_collisions`  
WHERE  
contributing_factor_vehicle_1 IS NULL
```

```
SELECT  
COUNT(*)  
FROM `bigquery-public-  
data.new_york_mv_collisions.nypd_mv_collisions`  
WHERE  
contributing_factor_vehicle_2 IS NULL
```

11 Introducing GROUP BY

We will work with the table: `usa_names.usa_1910_2013`

The GROUP BY clause groups a set of rows into a set of summary rows by values of columns or expressions. The GROUP BY clause returns one row for each group. In other words, it reduces the number of rows in the result set.

To get the total count of males and females:

```
SELECT  
    gender,  
    COUNT(gender) AS gender_count  
FROM `bigquery-public-data.usa_names.usa_1910_2013`  
GROUP BY  
    gender
```

To get the total count of people with a distinct name

```
SELECT  
    name,  
    COUNT(name) AS name_count  
FROM `bigquery-public-data.usa_names.usa_1910_2013`  
GROUP BY  
    name
```

To get the total count of people from a distinct state

```
SELECT
    state,
    COUNT(state) AS state_count
FROM
    `bigquery-public-data.usa_names.usa_1910_2013`
GROUP BY
    state
```

To filter the results returned from the count, we use the HAVING clause

```
SELECT
    state,
    COUNT(state) AS state_count
FROM
    `bigquery-public-data.usa_names.usa_1910_2013`
GROUP BY
    state
HAVING
    state_count > 100000
```

To filter and sort the results returned from the count we can use the ORDER BY clause

```
SELECT
    state,
    COUNT(state) AS state_count
FROM
    `bigquery-public-data.usa_names.usa_1910_2013`
GROUP BY
    state
HAVING
    state_count > 100000
ORDER BY
    state_count DESC
```

To filter before grouping, we can use the WHERE clause which comes before the GROUP BY. This gets applied first to the table to filter it: (for e.g. finding the count of all females in all the states)

```
SELECT
    state,
    COUNT(state) AS state_count
FROM
    `bigquery-public-data.usa_names.usa_1910_2013`
WHERE
    gender = 'F'
GROUP BY
    state
ORDER BY
    state_count DESC
```


12 Using GROUP BY with aggregate functions

We will work with the table: `chicago_taxi_trips.taxi_trips`

Typically, the MAX, MIN, AVG, SUM, and COUNT functions are applied along with the GROUP BY statement to perform aggregation over the grouped records.

These aggregate functions can be applied individually:

```
SELECT
  payment_type,
  MIN(trip_total) AS min_trip_total
FROM
  `bigquery-public-data.chicago_taxi_trips.taxi_trips`
GROUP BY
  payment_type
ORDER BY
  payment_type
```

```
SELECT
  payment_type,
  MAX(trip_total) AS min_trip_total
FROM
  `bigquery-public-data.chicago_taxi_trips.taxi_trips`
GROUP BY
  payment_type
ORDER BY
  payment_type
```

They can also be combined into a single query:

```
SELECT
  payment_type,
  COUNT(DISTINCT unique_key) AS num_trips,
  SUM(trip_total) AS sum_trip_total,
  AVG(trip_total) AS avg_trip_total,
  MAX(trip_total) AS max_trip_total,
  MIN(trip_total) AS min_trip_total
FROM
  `bigquery-public-data.chicago_taxi_trips.taxi_trips`
GROUP BY
  payment_type
ORDER BY
  payment_type
```

You can also add in the WHERE clause to filter before performing the GROUP by and applying the aggregation function:

```

SELECT
  payment_type,
  COUNT(DISTINCT unique_key) AS num_trips,
  SUM(trip_total) AS sum_trip_total,
  AVG(trip_total) AS avg_trip_total,
  MAX(trip_total) AS max_trip_total,
  MIN(trip_total) AS min_trip_total
FROM
  `bigquery-public-data.chicago_taxi_trips.taxi_trips`
WHERE
  payment_type IN ('Cash', 'Credit Card', 'Mobile')
GROUP BY
  payment_type
ORDER BY
  num_trips DESC

```

13 Performing JOINS

Joining tables means to JOIN columns from one table onto another table given some join condition/criteria. This is typically done when there is a column or set of columns in common between tables.

Tables to use:

```

census_bureau_international.midyear_population
census_bureau_international.country_names_area

```

Query to use:

```

SELECT
  m.year,
  m.country_name AS country,
  m.midyear_population AS population,
  a.country_area AS area
FROM
  `bigquery-public-
data.census_bureau_international.midyear_population` m
LEFT JOIN
  `bigquery-public-
data.census_bureau_international.country_names_area` a
ON
  m.country_code = a.country_code
ORDER BY
  year,
  country

```

14 Looker Studio background

General Overview of Looker Studio and its use in BI

<https://cloud.google.com/looker-studio>

You can browse through and experiment with some of the templates available
<https://lookerstudio.google.com/navigation/templates>